

Assignment 1

CS-351

Fall 2021

Due Date: 11/5/2021 at 11:59 pm (No extensions)

You may work in groups of 1-5 (only one person from the group should submit)

You may submit as many times as you like

Canvas will save only the last attempt

Goals:

1. To gain hands-on experience with `fork()`, `exec()`, and `wait()` system calls.
2. To master the basics of multi-process application development.
3. To appreciate the performance and fault-tolerance benefits of multi-process applications.
4. To implement a multi-process downloader application.

Part I Overview

Write a simple shell program called `shell`. When ran, the parent process will print out the prompt which looks like:

```
cmd>
```

and will then wait for the user to enter a command such as `ls`, `ps`, or any other command. For example:

```
cmd> ls
```

The parent process will then use `fork()` in order to create a child and then call `wait()` in order to wait for the child to terminate. The child process will use `execlp()` in order to replace its program with the program entered at the command line. After the child produces the output, the parent will prompt the user for another command. For example:

```
cmd>ls
file.txt fork fork.c mystery1 mystery1.cpp mystery2 shell shell.cpp
```

The parent will repeat the above sequence until the user types the `exit` command, which will cause the parent process to exit.

Sample/Skeleton Codes

A basic forking example (`fork.cpp`) had been provided in the class notes as well as in the zip bundle of sample files/skeletons accompanying this assignment called `samples.zip`. The bundle also includes a skeleton file for the shell with parts to complete marked with `TODO`. You are **not required** to use these files, but may find them helpful.

NOTE: Please make sure to error-check all system calls in this assignment. This is very important in practice and can also save you hours of debugging frustration. `fork()`, `execlp()`, and `wait()` will return -1 on error. Hence, you need to always check their return values and terminate your program if the return value is -1. For example:

```
pid_t pid = fork()

if(pid < 0)
{
    perror('fork');
    exit(-1);
}
```

The `perror()` function above will print out `fork` followed by the explanation of the error.

Part II Overview

File downloaders are programs used for downloading files from the Internet. In this assignment you will implement two different types of *multi-process downloaders* (i.e., file downloaders that comprise multiple processes):

1. **a serial file downloader** which downloads files one by one.
2. **a parallel file downloader** which downloads multiple files in parallel.

You will then compare the performance of the two types of downloaders. Both downloaders will use the Linux `wget` program in order to perform the actual downloading. The usage of the `wget` is simple: `wget <FILE URL>`. For example, running from command line the following command:

```
wget http://old-releases.ubuntu.com/releases/15.04/ubuntu-15.04-desktop-amd64.iso
```

will download the Ubuntu Linux iso image to the current directory. Before proceeding with the assignment, you may want to take a moment to experiment with the `wget` command.

I can also run the program as:

```
wget http://emerald.ecs.fullerton.edu/file1 -o log.txt
```

which will download the specified file and save the log information (which in the previous example would be printed on the terminal) to file `log.txt`.

In your program, the parent process shall first read the file, `urls.txt`, containing the URLs of the files to be downloaded (there can be any number of URLs in the file). `urls.txt` shall have

the following format:

```
<URL1>
<URL2>
.
.
.
<URLN>
```

For example:

```
http://emerald.ecs.fullerton.edu/file1
http://emerald.ecs.fullerton.edu/file2
http://emerald.ecs.fullerton.edu/file3
```

Next, the parent process shall fork the child processes. Each created child process shall use the `execlp()` system call to replace its executable image with that of the `wget` program. The two types downloaders are described in detail below.

The two downloaders shall be implemented as **separate programs**. The serial downloader program shall be called `serial.c` (or `.cpp` extension if you use C++). The parallel downloader program shall be called `parallel.c` (or `.cpp` extension if you use C++).

Serial Downloader

The serial downloader shall download files one by one. After the parent process has read and parsed the `urls.txt` file, it shall proceed as follows:

1. The parent process forks off a child process.
2. The child uses `execlp("/usr/bin/wget", "wget", <URL STRING1>, ... other arguments NULL)` system call in order to replace its program with `wget` program that will download the first file in `urls.txt` (i.e. the file at URL `<URL STRING1>`). Also, the child should save the logging information produced by `wget` into a file with a unique name. Specifically, `log1.txt` for URL1, `log2.txt` for URL2, etc. `wget` can be instructed to save e.g., the log information from downloading URL1 to `log1.txt` as `wget http://www.myurl1.com -o log1.txt`, which will download file at `http://www.myurl1.com` and save the logging information into `log1.txt` file. To achieve this effect, you will need to specify the proper arguments to `execlp()` as we discussed in class. Also, files numbered `log1.txt...logn.txt` can be generated using an approach similar to the one provided in `numberedfilegen.cpp`.
3. The parent executes a `wait()` system call until the child exits.
4. The parent forks off another child process which downloads the next file specified in `urls.txt`.
5. Repeat the above steps until all files are downloaded.

Parallel Downloader

1. The parent forks off n children, where n is the number of URLs in `urls.txt`.
2. Each child executes `execlp("/usr/bin/wget", "wget", <URL STRING>, NULL)` system call where each `<URL STRING>` is a distinct URL in `urls.txt`.
3. The parent calls `wait()` (n times in a row) and waits for all children to terminate. *Same requirements for saving the logging information apply as for the serial downloader.*
4. The parent exits.

Please note:

- While the parallel downloader executes, the outputs from different children may inter-mingle. This is acceptable.
- `fork.c` file posted on Canvas provides an example of using `fork()`, `execlp()`, and `wait()` system calls. Please feel free to modify it in order to complete the above tasks.

Performance Comparison

Use the `time` program to measure the execution time for the two downloaders. For example:

```
time ./serial
real 0m10.009s
user 0m0.008s
sys 0m0.000s
```

The column titled `real` gives the execution time in seconds. Please get the execution times for both downloaders using the following `urls.txt` file:

```
http://emerald.ecs.fullerton.edu/file1
http://emerald.ecs.fullerton.edu/file2
http://emerald.ecs.fullerton.edu/file3
```

Your execution times should be submitted along with your code (see the section titled “Submission Guidelines”).

In your submission, please include the answers to the following questions (you may need to do some research):

1. In the output of `time`, what is the difference between `real`, `user`, and `sys` times?
2. Which is longer: `user` time or `sys` time? Use your knowledge to explain why.
3. When downloading the files above, which downloader finishes faster? Why do you think that is?
4. Repeat the experiment for 10 files (any reasonably large-sized files, e.g., 100 MB, will do). Is the downloader in the previous question still faster? If not so, why do you think that is?

BONUS

Implement a multi-process linear search.

The parent process shall load a file of strings into an array (or a vector), split the array into n sections, and then fork off n children. Each child process shall search one of the n sections of the array, for the specified string. If the child does not find the string, it terminates with the exit code of 1. Otherwise, it terminates with the exit code of 0.

The parent, meanwhile, continuously executes `wait()`. Whenever, one of the child processes terminates, the `wait()` will unblock, and the parent process shall check the exit code of the child. The exit status of the child can be checked using the `WEXITSTATUS` macro as follows:

```
//The variable to hold the exit status int exit_status;

//Other code ....

if(wait(&exit_status) < 1) { perror("wait"); exit(-1); }
if(WEXITSTATUS(exit_status) == 0) { ... } ...
```

If all children terminate with the code of 1, then the parent prints "No string found" to the terminal. Otherwise, the parent terminates all child processes and exits. A child process can be terminated using the `kill()` system call. For example, `kill(1234,SIGKILL)` will terminate the process with process id 1234 (be sure to include the header files `#include <sys/types.h>` and `#include <signal.h>`; otherwise your program may not compile).

Technical Details

The program shall be ran using the following command line:

```
./multi-search <FILE NAME> <KEY> <NUMBER OF PROCESSES>
```

Where `<FILE NAME>` is the name of the file containing the strings, `<NUMBER OF PROCESSES>` is the number of child processes, and `<KEY>` is the string to search for. For example, `./multi-search strings.txt abcd 10` tells the program to split the task of searching for string `abcd` in file `string.txt` amongst 10 child processes.

SUBMISSION GUIDELINES:

- ***This assignment MUST be completed using C or C++ on Linux.*** item You may work in groups of 4.
- ***Your assignment must compile and run on Tuffix.*** Please contact the CS office if you have trouble logging in.
- Please hand in your source code electronically (do not submit .o or executable code) through **CANVAS**.

- You must make sure that the code compiles and runs correctly.
- Write a README file (text file, do not submit a .doc file) which contains
 - Your name and email address.
 - The programming language you used (i.e., C or C++).
 - How to execute your program.
 - The execution times for both downloaders.
 - The answers to all questions above.
 - Whether you implemented the extra credit.
 - Anything special about your submission that we should take note of.
- Place all your files under one directory with a unique name (such as p1-[userid] for assignment 1, e.g. p1-mgofman1).
- Tar the contents of this directory using the following command. `tar cvf [directory_name].tar [directory_name]` E.g. `tar -cvf p1-mgofman1.tar p1-mgofman1/`
- Use CANVAS to upload the tared file you created above.

Grading guideline:

- All programs compile: 5'
- Correctly shell: 25'
- Correct serial downloader: 30'
- Correct parallel downloader: 30'
- Execution times for both downloaders: 5'
- README file: 5'
- Bonus: 15'
- Late submissions shall be penalized 10%. No assignments shall be accepted after 24 hours.

Academic Honesty:

Academic Honesty: All forms of cheating shall be treated with utmost seriousness. You may discuss the problems with other students, however, you must write your **OWN codes and solutions**. Discussing solutions to the problem is **NOT** acceptable (unless specified otherwise). Copying an assignment from another student or allowing another student to copy your work **may lead to an automatic F for this course**. Moss shall be used to detect plagiarism in programming assignments. If you have any questions about whether an act of collaboration may be treated as academic dishonesty, please consult the instructor before you collaborate. Details posted at <http://www.fullerton.edu/senate/documents/PDF/300/UPS300-021.pdf>.